

IF2130 – Organisasi dan Arsitektur Komputer

sumber: Greg Kesden, CMU 15-213, 2012

Cache

Achmad Imam Kistijantoro (imam@staff.stei.itb.ac.id)

Rahmat Mulyawan

Infall Syafalni

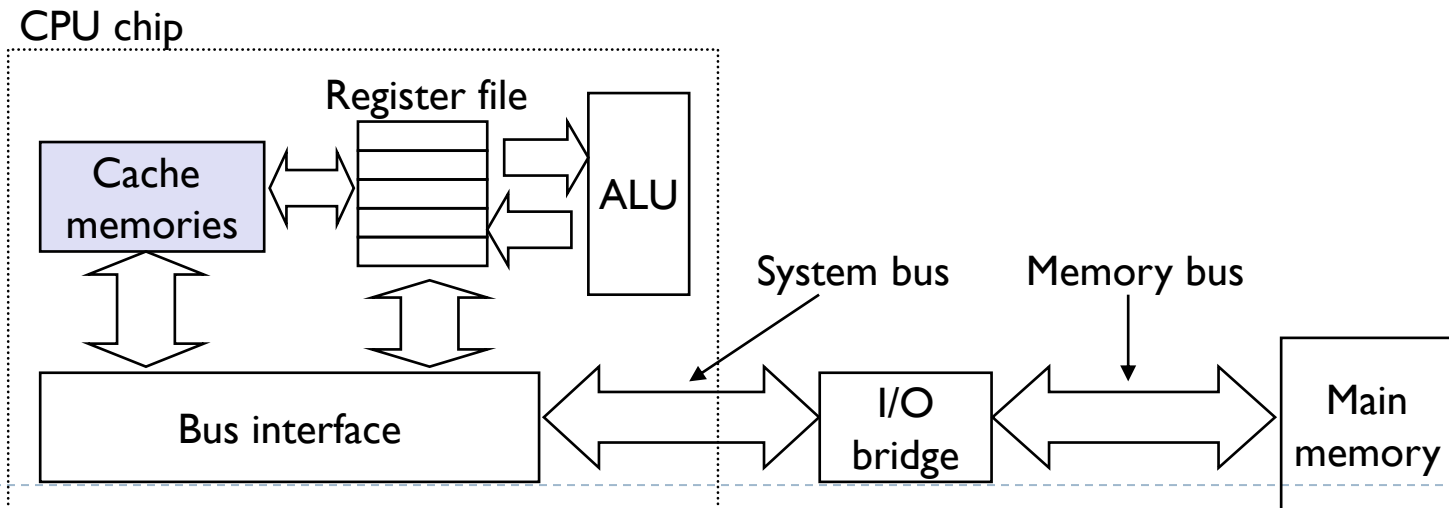
Today

- ▶ **Cache memory organization and operation**
- ▶ Performance impact of caches
 - ▶ The memory mountain
 - ▶ Rearranging loops to improve spatial locality
 - ▶ Using blocking to improve temporal locality




Cache Memories

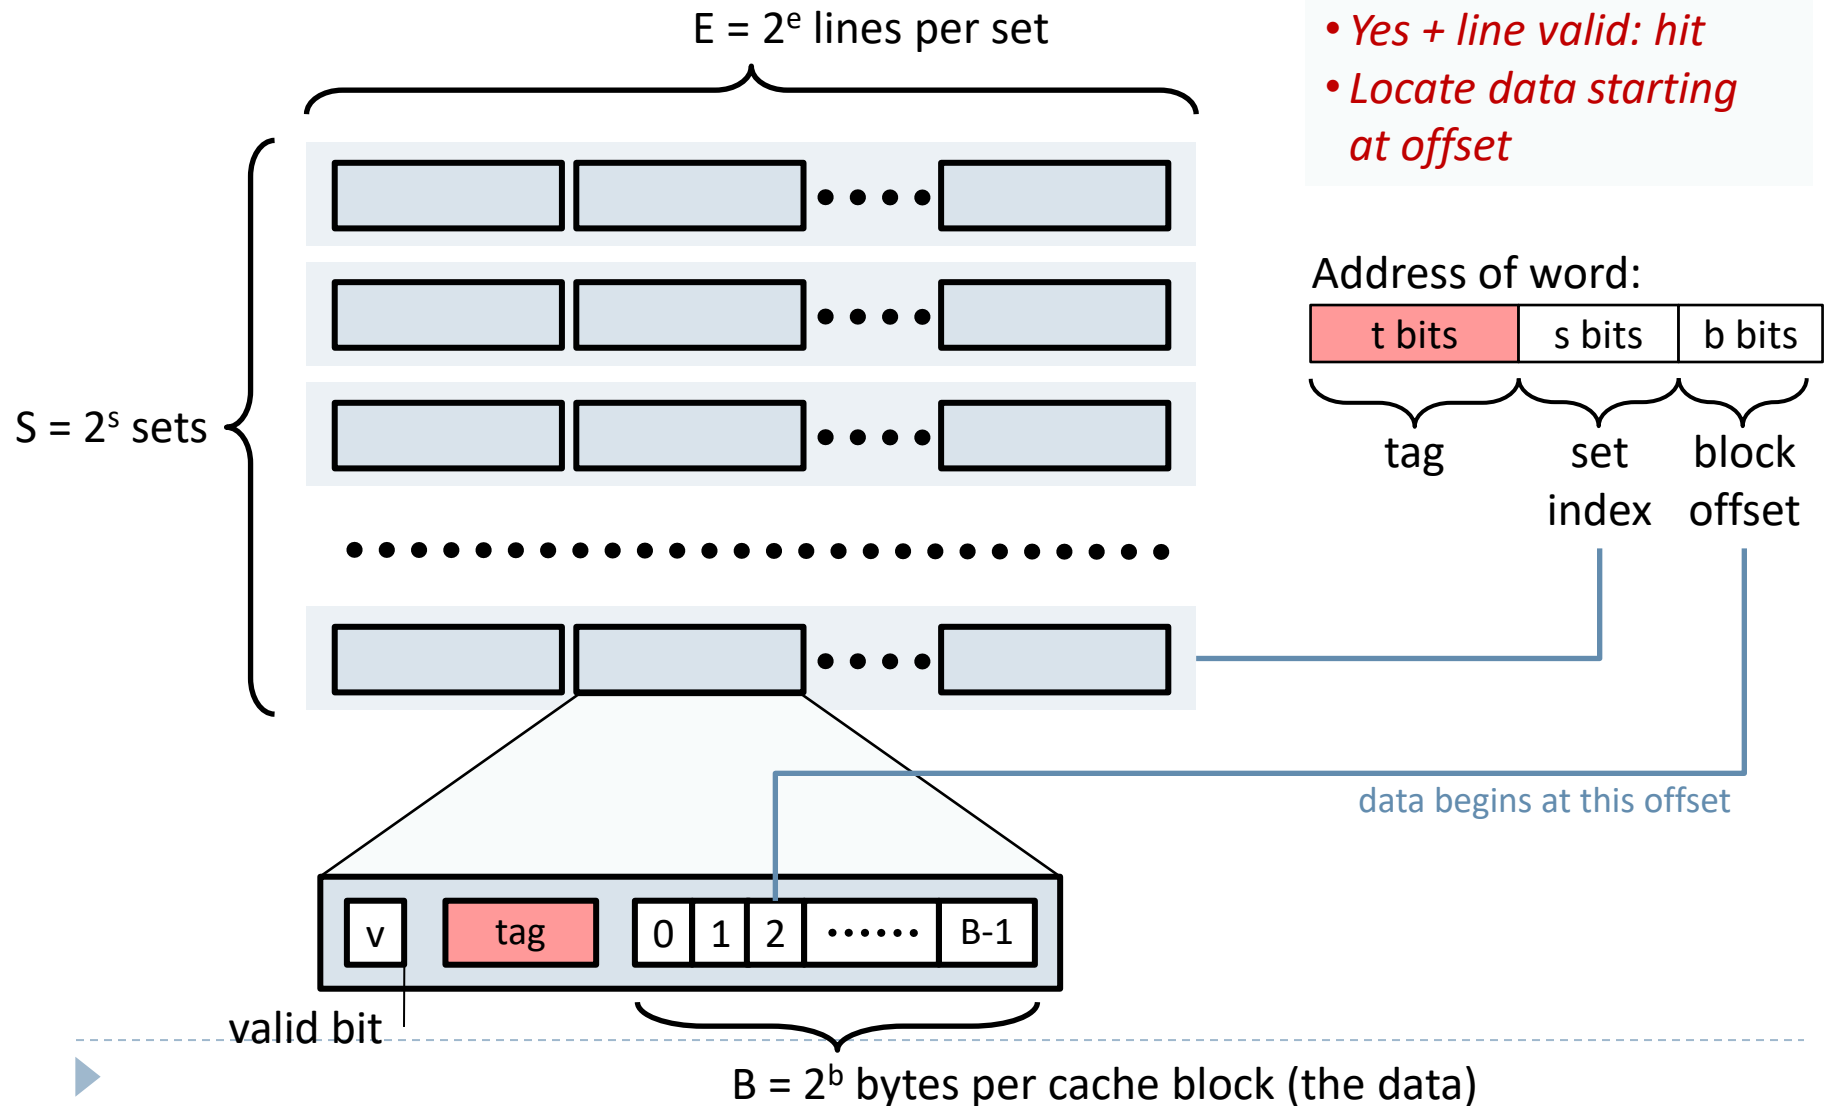
- ▶ **Cache memories** are small, fast SRAM-based memories managed automatically in hardware.
 - ▶ Hold frequently accessed blocks of main memory
- ▶ CPU looks first for data in caches (e.g., L1, L2, and L3), then in main memory.
- ▶ Typical system structure:



General Cache Organization (S, E, B)



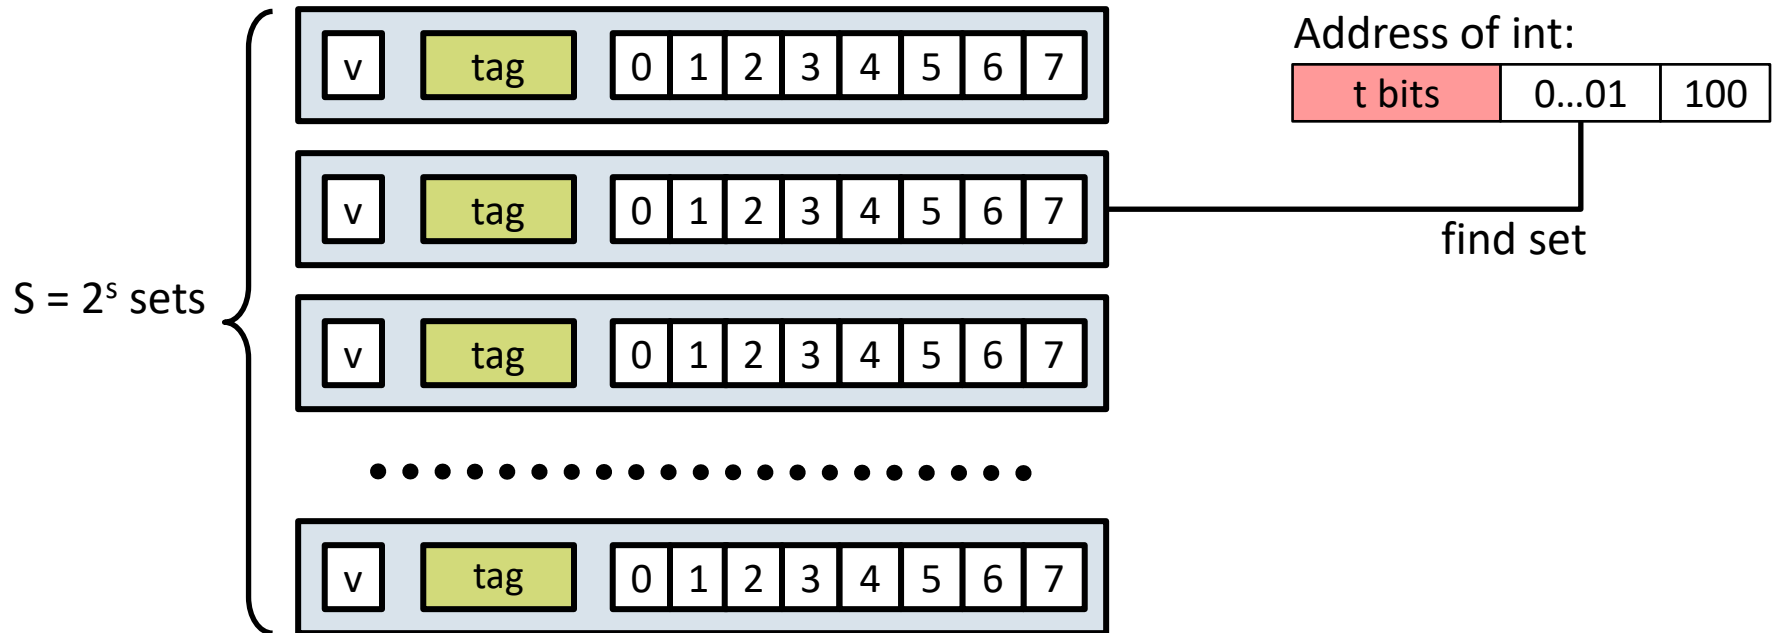
Cache Read



Example: Direct Mapped Cache ($E = 1$)

Direct mapped: One line per set


Assume: cache block size 8 bytes



Example: Direct Mapped Cache (E = 1)

Direct mapped: One line per set


Assume: cache block size 8 bytes



Example: Direct Mapped Cache (E = 1)

Direct mapped: One line per set

Assume: cache block size 8 bytes



No match: old line is evicted and replaced

Direct-Mapped Cache Simulation

t=1	s=2	b=1
x	xx	x

M=16 byte addresses, B=2 bytes/block,
S=4 sets, E=1 Blocks/set

Address trace (reads, one byte per read):

0	[<u>0000</u> ₂],	miss
1	[<u>0001</u> ₂],	hit
7	[<u>0111</u> ₂],	miss
8	[<u>1000</u> ₂],	miss
0	[<u>0000</u> ₂]	miss

	v	Tag	Block
Set 0	1	0	M[0-1]
Set 1			
Set 2			
Set 3	1	0	M[6-7]



A Higher Level Example

```
int sum_array_rows(double a[16][16])
{
    int i, j;
    double sum = 0;

    for (i = 0; i < 16; i++)
        for (j = 0; j < 16; j++)
            sum += a[i][j];
    return sum;
}
```

```
int sum_array_cols(double a[16][16])
{
    int i, j;
    double sum = 0;

    for (j = 0; j < 16; j++)
        for (i = 0; i < 16; i++)
            sum += a[i][j];
    return sum;
}
```

Ignore the variables sum, i, j

assume: cold (empty) cache,
a[0][0] goes here



32 B = 4 doubles

blackboard

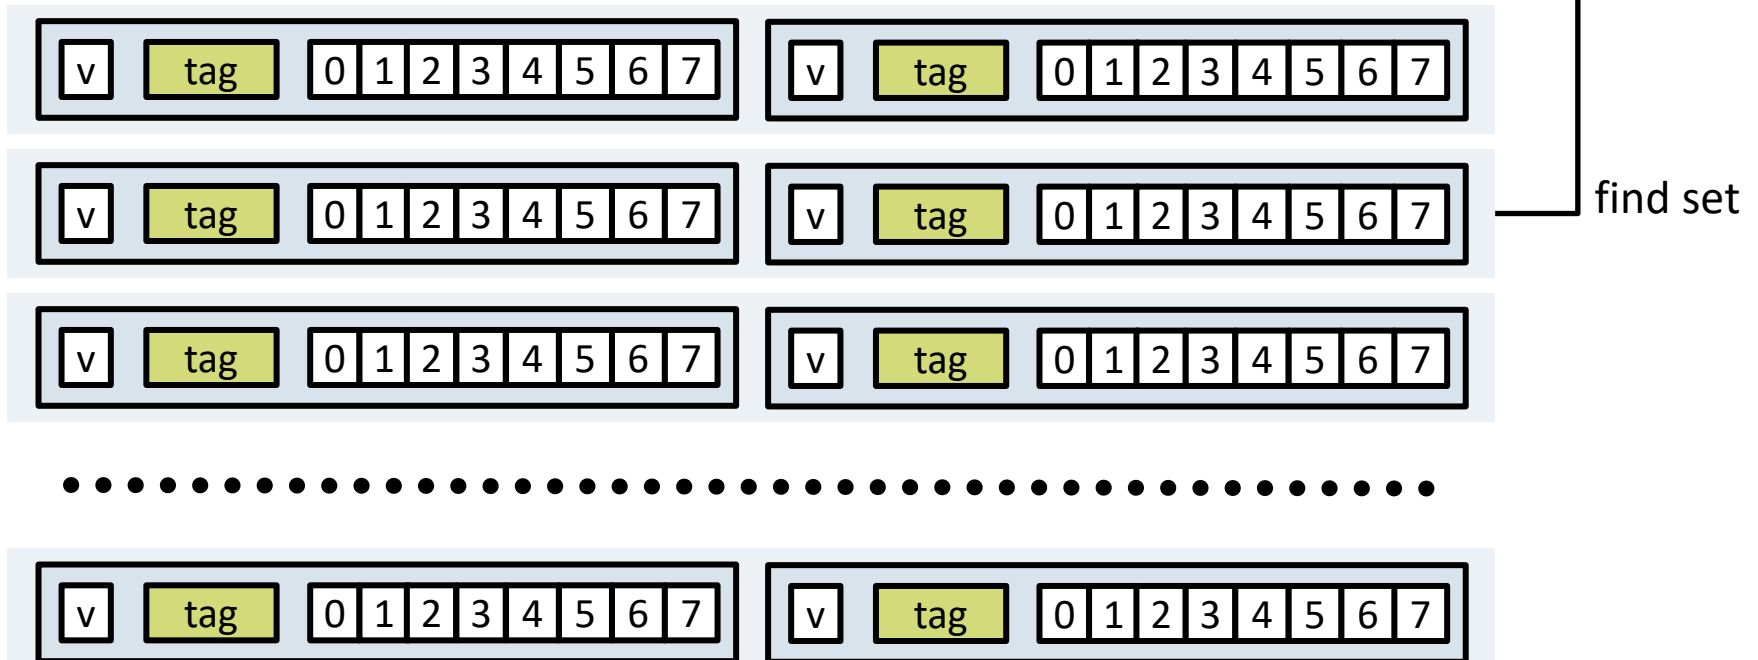
E-way Set Associative Cache (Here: E = 2)

E = 2: Two lines per set

Assume: cache block size 8 bytes

Address of short int:

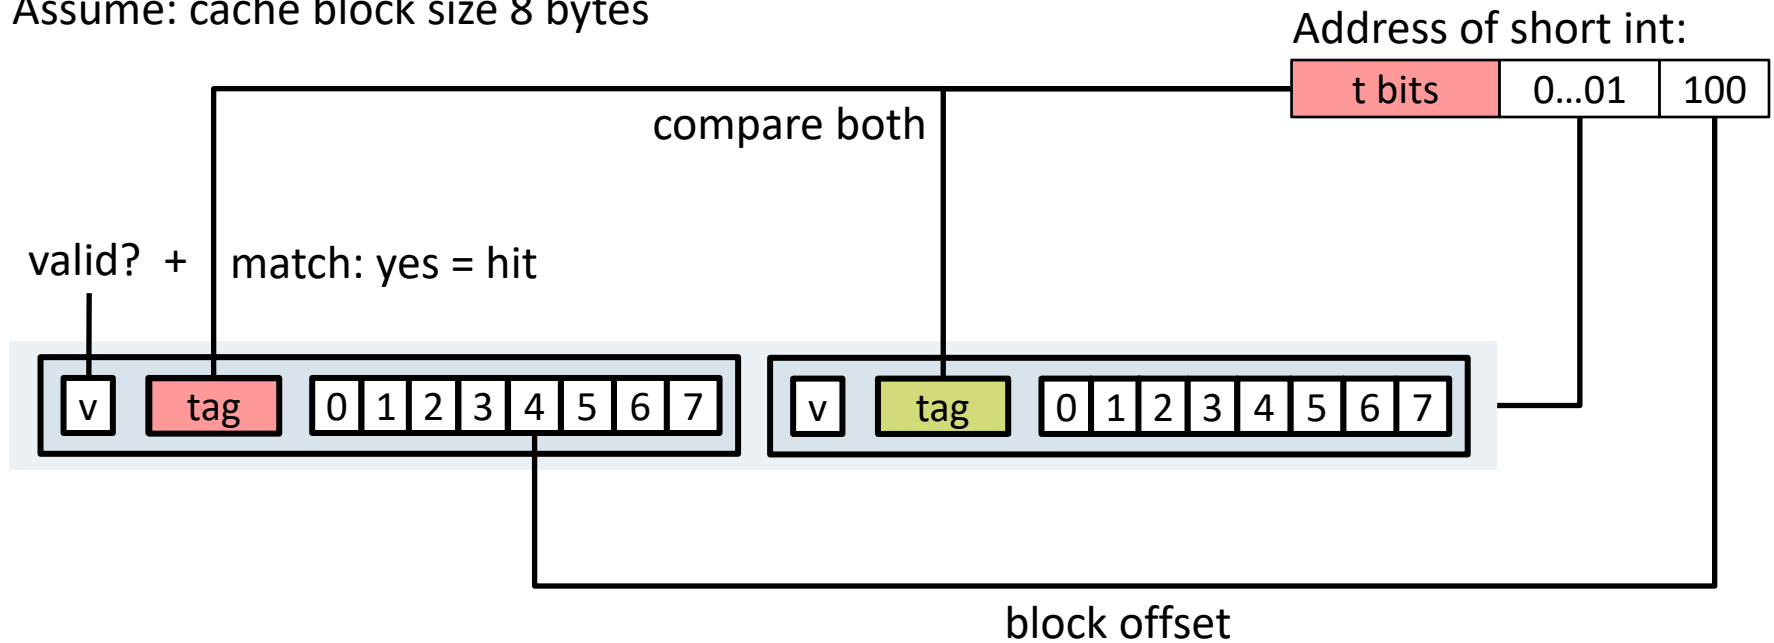
t bits	0...01	100
--------	--------	-----



E-way Set Associative Cache (Here: E = 2)

E = 2: Two lines per set

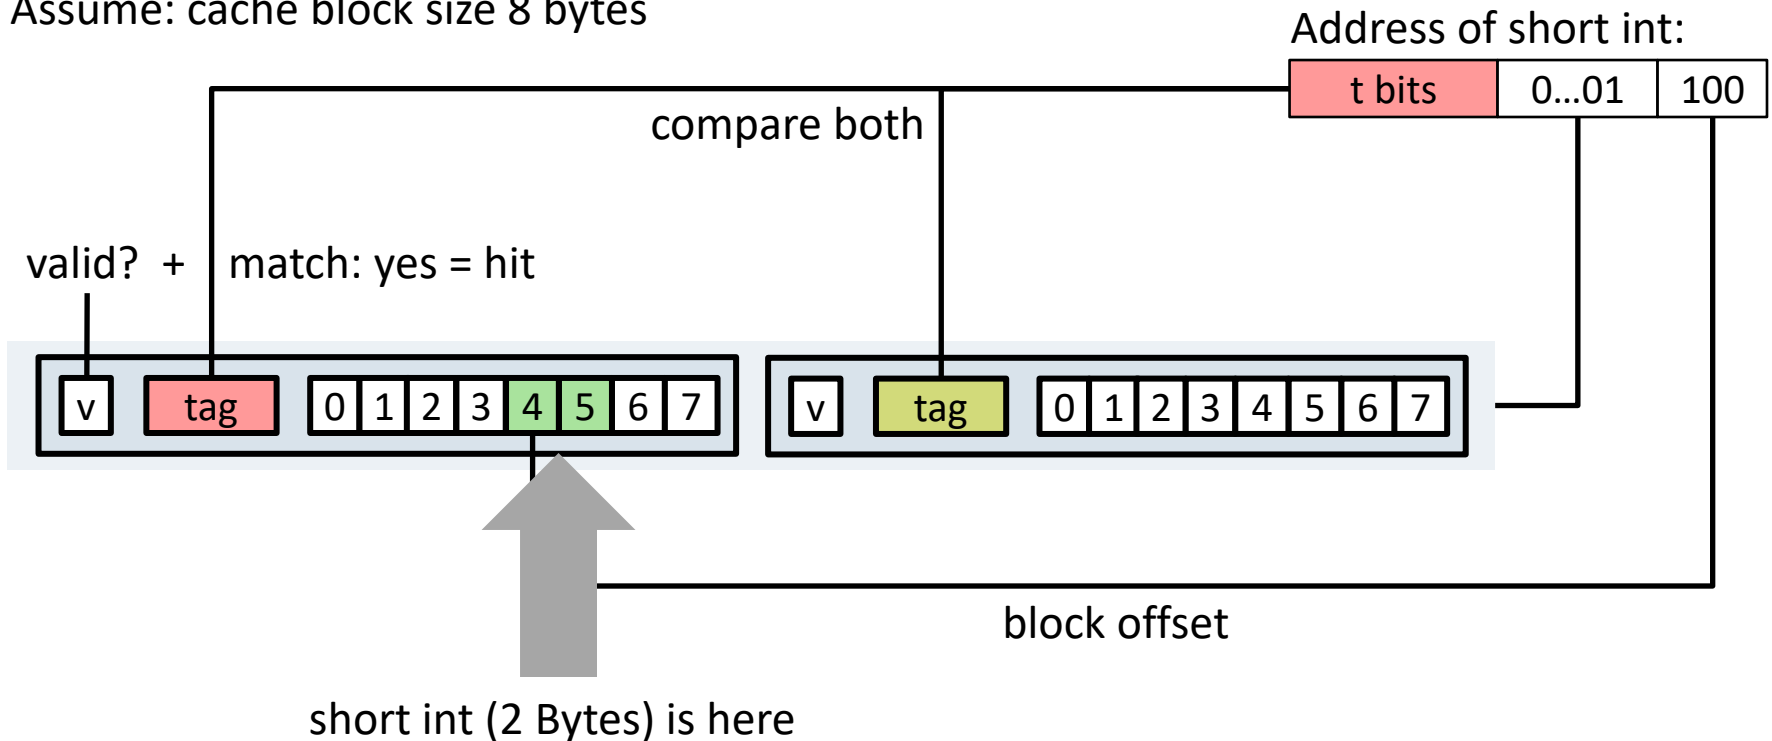
Assume: cache block size 8 bytes



E-way Set Associative Cache (Here: E = 2)

E = 2: Two lines per set

Assume: cache block size 8 bytes



No match:

- One line in set is selected for eviction and replacement
- Replacement policies: random, least recently used (LRU), ...

2-Way Set Associative Cache Simulation

t=2	s=1	b=1
xx	x	x

M=16 byte addresses, B=2 bytes/block,
S=2 sets, E=2 blocks/set

Address trace (reads, one byte per read):

0	[0000 ₂],	miss
1	[0001 ₂],	hit
7	[0111 ₂],	miss
8	[1000 ₂],	miss
0	[0000 ₂]	hit

	v	Tag	Block
Set 0	1	00	M[0-1]
	1	10	M[8-9]
Set 1	1	01	M[6-7]
	0		

A Higher Level Example

Ignore the variables sum, i, j


```
int sum_array_rows(double a[16][16])
{
    int i, j;
    double sum = 0;

    for (i = 0; i < 16; i++)
        for (j = 0; j < 16; j++)
            sum += a[i][j];
    return sum;
}
```

```
int sum_array_cols(double a[16][16])
{
    int i, j;
    double sum = 0;

    for (j = 0; j < 16; j++)
        for (i = 0; i < 16; i++)
            sum += a[i][j];
    return sum;
}
```

assume: cold (empty) cache,
a[0][0] goes here



32 B = 4 doubles

blackboard

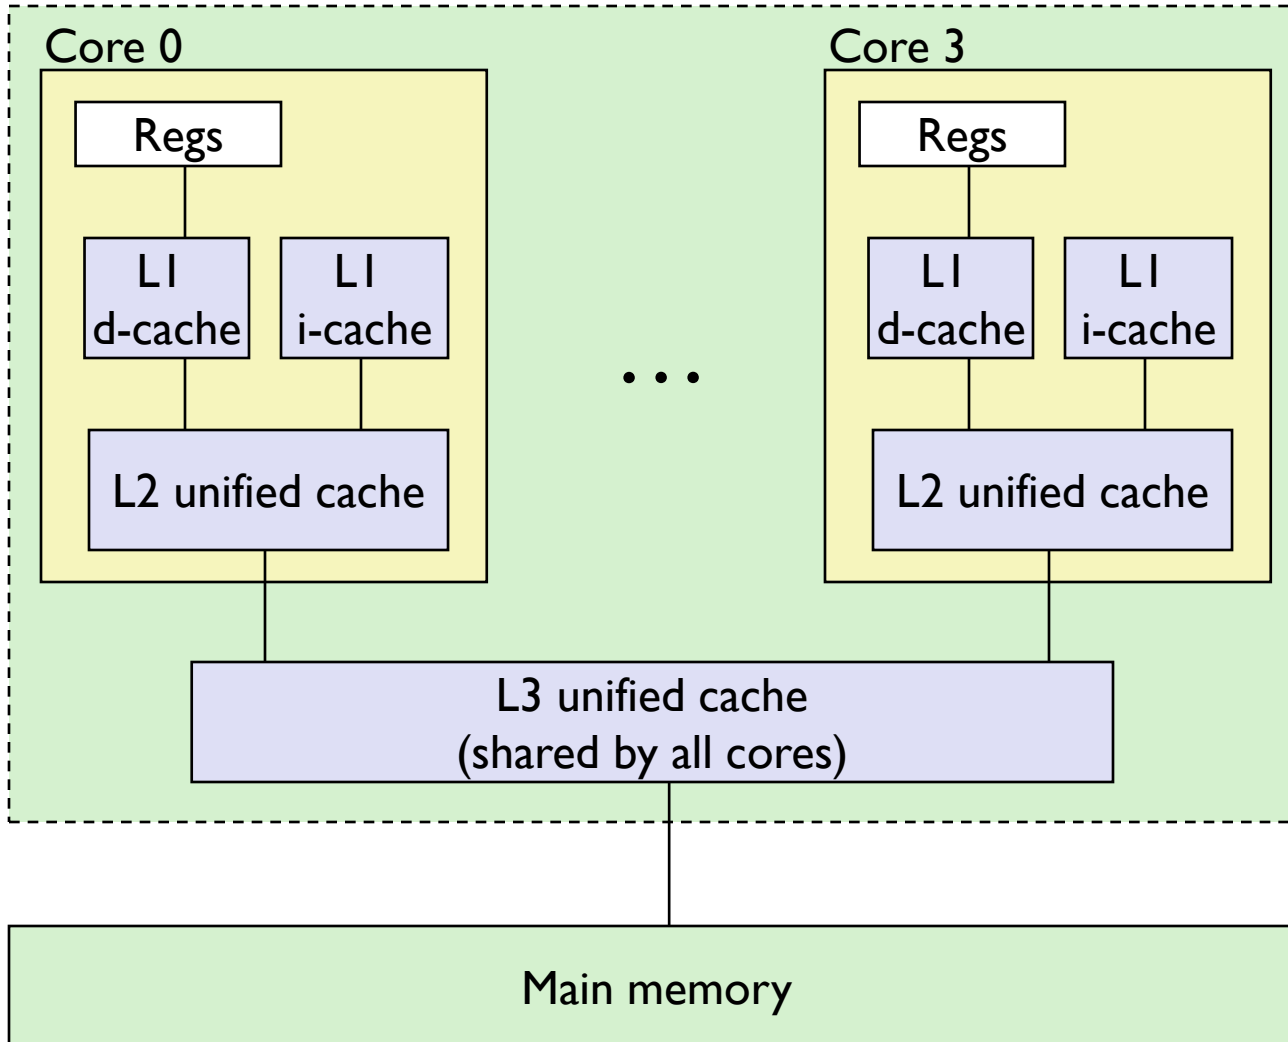
What about writes?

- ▶ Multiple copies of data exist:
 - ▶ L1, L2, Main Memory, Disk
- ▶ What to do on a write-hit?
 - ▶ **Write-through** (write immediately to memory)
 - ▶ **Write-back** (defer write to memory until replacement of line)
 - ▶ Need a dirty bit (line different from memory or not)
- ▶ What to do on a write-miss?
 - ▶ **Write-allocate** (load into cache, update line in cache)
 - ▶ Good if more writes to the location follow
 - ▶ **No-write-allocate** (writes immediately to memory)
- ▶ Typical
 - ▶ Write-through + No-write-allocate
 - ▶ **Write-back + Write-allocate**



Intel Core i7 Cache Hierarchy

Processor package



L1 i-cache and d-cache:
32 KB, 8-way,
Access: 4 cycles

L2 unified cache:
256 KB, 8-way,
Access: 11 cycles

L3 unified cache:
8 MB, 16-way,
Access: 30-40 cycles

Block size: 64 bytes for
all caches.

Cache Performance Metrics

▶ Miss Rate

- ▶ Fraction of memory references not found in cache (misses / accesses)
= $1 - \text{hit rate}$
- ▶ Typical numbers (in percentages):
 - ▶ 3-10% for L1
 - ▶ can be quite small (e.g., $< 1\%$) for L2, depending on size, etc.

▶ Hit Time

- ▶ Time to deliver a line in the cache to the processor
 - ▶ includes time to determine whether the line is in the cache
- ▶ Typical numbers:
 - ▶ 1-2 clock cycle for L1
 - ▶ 5-20 clock cycles for L2

▶ Miss Penalty

- ▶ Additional time required because of a miss
 - ▶ typically 50-200 cycles for main memory (Trend: increasing!)



Lets think about those numbers

- ▶ Huge difference between a hit and a miss
 - ▶ Could be 100x, if just L1 and main memory
- ▶ Would you believe 99% hits is twice as good as 97%?
 - ▶ Consider:
cache hit time of 1 cycle
miss penalty of 100 cycles
 - ▶ Average access time:
97% hits: $1 \text{ cycle} + 0.03 * 100 \text{ cycles} = \mathbf{4 \text{ cycles}}$
99% hits: $1 \text{ cycle} + 0.01 * 100 \text{ cycles} = \mathbf{2 \text{ cycles}}$
- ▶ This is why “miss rate” is used instead of “hit rate”



Writing Cache Friendly Code

- ▶ Make the common case go fast
 - ▶ Focus on the inner loops of the core functions
- ▶ Minimize the misses in the inner loops
 - ▶ Repeated references to variables are good (**temporal locality**)
 - ▶ Stride-1 reference patterns are good (**spatial locality**)

Key idea: Our qualitative notion of locality is quantified through our understanding of cache memories.



Today

- ▶ Cache organization and operation
- ▶ **Performance impact of caches**
 - ▶ The memory mountain
 - ▶ Rearranging loops to improve spatial locality
 - ▶ Using blocking to improve temporal locality



The Memory Mountain

- ▶ **Read throughput** (read bandwidth)
 - ▶ Number of bytes read from memory per second (MB/s)
- ▶ **Memory mountain:** Measured read throughput as a function of spatial and temporal locality.
 - ▶ Compact way to characterize memory system performance.



Memory Mountain Test Function

```
long data[MAXELEMS]; /* Global array to traverse */
```

```
/* test - Iterate over first "elems" elements of  
 *   array "data" with stride of "stride", using  
 *   using 4x4 loop unrolling.  
 */
```

```
int test(int elems, int stride) {  
    long i, sx2=stride*2, sx3=stride*3, sx4=stride*4;  
    long acc0 = 0, acc1 = 0, acc2 = 0, acc3 = 0;  
    long length = elems, limit = length - sx4;
```

```
/* Combine 4 elements at a time */
```

```
for (i = 0; i < limit; i += sx4) {  
    acc0 = acc0 + data[i];  
    acc1 = acc1 + data[i+stride];  
    acc2 = acc2 + data[i+sx2];  
    acc3 = acc3 + data[i+sx3];  
}
```

```
/* Finish any remaining elements */
```

```
for (; i < length; i++) {  
    acc0 = acc0 + data[i];  
}  
return ((acc0 + acc1) + (acc2 + acc3));
```

```
}
```

mountain/mountain.c

Call `test()` with many combinations of `elems` and `stride`.

For each `elems` and `stride`:


1. Call `test()` once to warm up the caches.

2. Call `test()` again and measure the read throughput (MB/s)

The Memory Mountain

Core i7 Haswell
2.1 GHz
32 KB L1 d-cache
256 KB L2 cache
8 MB L3 cache
64 B block size

*Aggressive
prefetching*



Today

- ▶ Cache organization and operation
- ▶ Performance impact of caches
 - ▶ The memory mountain
 - ▶ Rearranging loops to improve spatial locality
 - ▶ Using blocking to improve temporal locality




Miss Rate Analysis for Matrix Multiply

► Assume:

- Line size = $32B$ (big enough for four 64-bit words)
- Matrix dimension (N) is very large
 - Approximate $1/N$ as 0.0
- Cache is not even big enough to hold multiple rows

► Analysis Method:

- Look at access pattern of inner loop



Matrix Multiplication Example

► Description:

- Multiply $N \times N$ matrices
- $O(N^3)$ total operations
- N reads per source element
- N values summed per destination
 - but may be able to hold in register

```
/* ijk */  
for (i=0; i<n; i++) {  
    for (j=0; j<n; j++) {  
        sum = 0.0;  
        for (k=0; k<n; k++)  
            sum += a[i][k] * b[k][j];  
        c[i][j] = sum;  
    }  
}
```

*Variable `sum`
held in register*

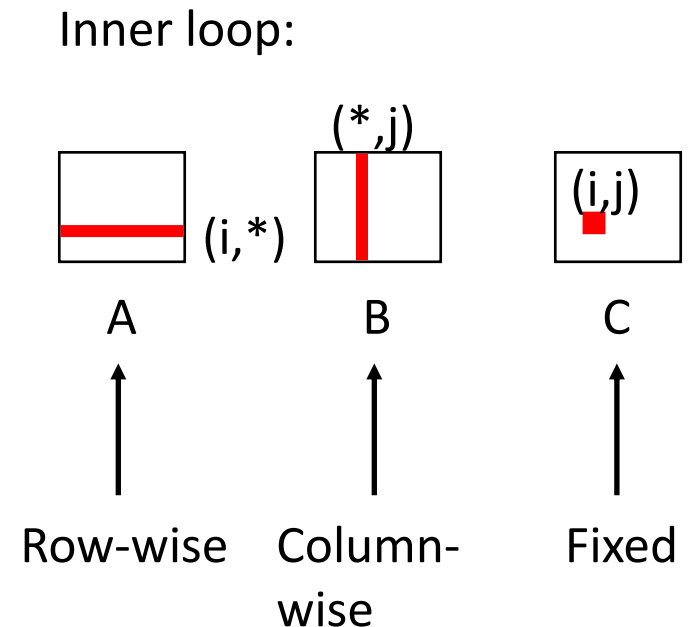
Layout of C Arrays in Memory (review)

- ▶ C arrays allocated in row-major order
 - ▶ each row in contiguous memory locations
- ▶ Stepping through columns in one row:
 - ▶ `for (i = 0; i < N; i++)`
 `sum += a[0][i];`
 - ▶ accesses successive elements
 - ▶ if block size (B) > 4 bytes, exploit spatial locality
 - ▶ compulsory miss rate = 4 bytes / B
- ▶ Stepping through rows in one column:
 - ▶ `for (i = 0; i < n; i++)`
 `sum += a[i][0];`
 - ▶ accesses distant elements
 - ▶ no spatial locality!
 - ▶ compulsory miss rate = 1 (i.e. 100%)



Matrix Multiplication (ijk)

```
/* ijk */
for (i=0; i<n; i++) {
  for (j=0; j<n; j++) {
    sum = 0.0;
    for (k=0; k<n; k++)
      sum += a[i][k] * b[k][j];
    c[i][j] = sum;
  }
}
```




Misses per inner loop iteration:

<u>A</u>	<u>B</u>	<u>C</u>
0.25	1.0	0.0

Matrix Multiplication (jik)

```
/* jik */
for (j=0; j<n; j++) {
  for (i=0; i<n; i++) {
    sum = 0.0;
    for (k=0; k<n; k++)
      sum += a[i][k] * b[k][j];
    c[i][j] = sum
  }
}
```

Inner loop:




Misses per inner loop iteration:

<u>A</u>	<u>B</u>	<u>C</u>
0.25	1.0	0.0

Matrix Multiplication (kij)

```
/* kij */
for (k=0; k<n; k++) {
    for (i=0; i<n; i++) {
        r = a[i][k];
        for (j=0; j<n; j++)
            c[i][j] += r * b[k][j];
    }
}
```

Inner loop:




Misses per inner loop iteration:

<u>A</u>	<u>B</u>	<u>C</u>
0.0	0.25	0.25

Matrix Multiplication (ikj)

```
/* ikj */
for (i=0; i<n; i++) {
    for (k=0; k<n; k++) {
        r = a[i][k];
        for (j=0; j<n; j++)
            c[i][j] += r * b[k][j];
    }
}
```

Inner loop:



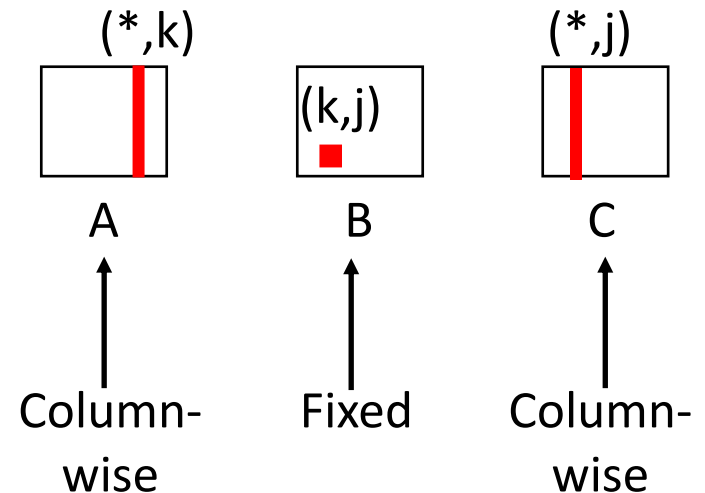
Misses per inner loop iteration:

<u>A</u>	<u>B</u>	<u>C</u>
0.0	0.25	0.25

Matrix Multiplication (jki)

```
/* jki */
for (j=0; j<n; j++) {
  for (k=0; k<n; k++) {
    r = b[k][j];
    for (i=0; i<n; i++)
      c[i][j] += a[i][k] * r;
  }
}
```

Inner loop:



Misses per inner loop iteration:

A
1.0


B
0.0

C
1.0

Matrix Multiplication (kji)

```
/* kji */  
for (k=0; k<n; k++) {  
    for (j=0; j<n; j++) {  
        r = b[k][j];  
        for (i=0; i<n; i++)  
            c[i][j] += a[i][k] * r;  
    }  
}
```

Inner loop:



Misses per inner loop iteration:

A
1.0

B
0.0

C
1.0

Summary of Matrix Multiplication

```
for (i=0; i<n; i++) {  
    for (j=0; j<n; j++) {  
        sum = 0.0;  
        for (k=0; k<n; k++)  
            sum += a[i][k] * b[k][j];  
        c[i][j] = sum;  
    }  
}
```

ijk (& jik):

- 2 loads, 0 stores
- misses/iter = 1.25

```
for (k=0; k<n; k++) {  
    for (i=0; i<n; i++) {  
        r = a[i][k];  
        for (j=0; j<n; j++)  
            c[i][j] += r * b[k][j];  
    }  
}
```

kij (& ikj):


- 2 loads, 1 store
- misses/iter = 0.5

```
for (j=0; j<n; j++) {  
    for (k=0; k<n; k++) {  
        r = b[k][j];  
        for (i=0; i<n; i++)  
            c[i][j] += a[i][k] * r;  
    }  
}
```

jki (& kji):

- 2 loads, 1 store
- misses/iter = 2.0

Core i7 Matrix Multiply Performance



Today


- ▶ Cache organization and operation
- ▶ Performance impact of caches
 - ▶ The memory mountain
 - ▶ Rearranging loops to improve spatial locality
 - ▶ Using blocking to improve temporal locality



Example: Matrix Multiplication

```
c = (double *) calloc(sizeof(double), n*n);

/* Multiply n x n matrices a and b */
void mmm(double *a, double *b, double *c, int n) {
    int i, j, k;
    for (i = 0; i < n; i++)
        for (j = 0; j < n; j++)
            for (k = 0; k < n; k++)
                c[i*n+j] += a[i*n + k]*b[k*n + j];
}
```



Cache Miss Analysis


▶ Assume:

- ▶ Matrix elements are doubles
- ▶ Cache block = 8 doubles
- ▶ Cache size $C \ll n$ (much smaller than n)

▶ First iteration:

- ▶ $n/8 + n = 9n/8$ misses

- ▶ Afterwards **in cache**:
(schematic)




Cache Miss Analysis

▶ Assume:

- ▶ Matrix elements are doubles
- ▶ Cache block = 8 doubles
- ▶ Cache size $C \ll n$ (much smaller than n)

▶ Second iteration:

- ▶ Again:
 $n/8 + n = 9n/8$ misses



▶ Total misses:


- ▶ $9n/8 * n^2 = (9/8) * n^3$



Blocked Matrix Multiplication


```
c = (double *) calloc(sizeof(double), n*n);

/* Multiply n x n matrices a and b */
void mmm(double *a, double *b, double *c, int n) {
    int i, j, k;
    for (i = 0; i < n; i+=B)
        for (j = 0; j < n; j+=B)
            for (k = 0; k < n; k+=B)
                /* B x B mini matrix multiplications */
                for (i1 = i; i1 < i+B; i++)
                    for (j1 = j; j1 < j+B; j++)
                        for (k1 = k; k1 < k+B; k++)
                            c[i1*n+j1] += a[i1*n + k1]*b[k1*n + j1];
}
```




Cache Miss Analysis

▶ Assume:

- ▶ Cache block = 8 doubles
- ▶ Cache size $C \ll n$ (much smaller than n)
- ▶ Three blocks  fit into cache: $3B^2 < C$


▶ First (block) iteration:

- ▶ $B^2/8$ misses for each block
- ▶ $2n/B * B^2/8 = nB/4$
(omitting matrix c)




Cache Miss Analysis

▶ Assume:

- ▶ Cache block = 8 doubles
- ▶ Cache size $C \ll n$ (much smaller than n)
- ▶ Three blocks  fit into cache: $3B^2 < C$

▶ Second (block) iteration:

- ▶ Same as first iteration
- ▶ $2n/B * B^2/8 = nB/4$



▶ Total misses:

- ▶ $nB/4 * (n/B)^2 = n^3/(4B)$

Summary

- ▶ No blocking: $(9/8) * n^3$
- ▶ Blocking: $1/(4B) * n^3$
- ▶ Suggest largest possible block size B , but limit $3B^2 < C$!
- ▶ Reason for dramatic difference:
 - ▶ Matrix multiplication has inherent temporal locality:
 - ▶ Input data: $3n^2$, computation $2n^3$
 - ▶ Every array elements used $O(n)$ times!
 - ▶ But program has to be written properly



Concluding Observations

- ▶ Programmer can optimize for cache performance
 - ▶ How data structures are organized
 - ▶ How data are accessed
 - ▶ Nested loop structure
 - ▶ Blocking is a general technique
- ▶ All systems favor “cache friendly code”
 - ▶ Getting absolute optimum performance is very platform specific
 - ▶ Cache sizes, line sizes, associativities, etc.
 - ▶ Can get most of the advantage with generic code
 - ▶ Keep working set reasonably small (temporal locality)
 - ▶ Use small strides (spatial locality)

